10 Characters and Strings

PURPOSE

1. To demonstrate the unique characteristics of character data
2. To view strings as an array of characters
3. To show how to input and output strings
4. To work with string functions

PROCEDURE

1. Students should read the Pre-lab Reading Assignment before coming to the lab.
2. Students should complete the Pre-lab Writing Assignment before coming to lab.
3. In the lab, students should complete labs assigned to them by the instructor.

Contents	Pre-requisites	Approximate completion time	Page number	Check when done
Pre-lab Reading Assignment		20 min.	176	
Pre-lab Writing Assignment	Pre-lab reading	10 min.	186	
LESSON 10A				
Lab 10.1 Character Testing and String Validation	Pre-lab reading	15 min.	187	
Lab 10.2 Case Conversion	Basic fundamental instructions	5 min.	190	
Lab 10.3 Using `getline() & get()`	Basic knowledge of character arrays	30 min.	192	
LESSON 10B				
Lab 10.4 String Functions—`strcat`	Basic knowledge of character arrays	15 min.	193	
Lab 10.5 Student Generated Code Assignments	Basic knowledge of character arrays	35 min.	193	

PRE-LAB READING ASSIGNMENT

Character Functions

C++ provides numerous *functions* for character testing. These functions will test a single character and return either a non-zero value (true) or zero (false). For example, isdigit tests a character to see if it is one of the digits between 0 and 9. So isdigit(7) returns a non-zero value whereas isdigit(y) and isdigit($) both return 0. We will not list all the character functions here. A complete list may be found in the text. The following program demonstrates some of the others. Note that the cctype header file must be included to use the character functions.

Sample Program 10.1:

```
// This program utilizes several functions for character testing

#include <iostream>
#include <cctype>
using namespace std;

int main()
{
    char input;

    cout << "Please Enter Any Character:" << endl;
    cin >> input;
    cout << "The character entered is " << input << endl << endl;
    cout << "The ASCII code for " << input << " is " << int(input)
        << endl;

    if (isalpha(input))     // tests to see if character is a letter
    {
      cout << "The character is a letter" << endl;

      if (islower(input))   // tests to see if letter is lower case
            cout << "The letter is lower case" << endl;

      if (isupper(input))   // tests to see if letter is upper case
            cout << "The letter is upper case" << endl;
    }

    else if (isdigit(input)) // tests to see if character is a digit
            cout << "The character you entered is a digit" << endl;

    else
            cout << "The character entered is not a letter nor a digit"
                    << endl;

    return 0;
}
```

In Lab 10.1 you will see a more practical application of character testing functions.

Character Case Conversion

The C++ library provides the `toupper` and `tolower` functions for converting the case of a character. `toupper` returns the uppercase equivalent for a letter and `tolower` returns the lower case equivalent. For example, `cout << tolower('F');` causes an `f` to be displayed on the screen. If the letter is already lowercase, then `tolower` will return the value unchanged. Likewise, any non-letter argument is returned unchanged by `tolower`. It should be clear to you now what `toupper` does to a given character.

While the `toupper` and `tolower` functions are conceptually quite simple, they may not appear to be very useful. However, the following program shows that they do have beneficial applications.

Sample Program 10.2:

```cpp
// This program shows how the toupper and tolower functions can be
// applied in a C++ program

#include <iostream>
#include <cctype>
#include <iomanip>
using namespace std;

int main()
{
    int week, total, dollars;
    float average;
    char choice;

    cout << showpoint << fixed << setprecision(2);

    do
    {
        total = 0;
        for(week = 1; week <= 4; week++)
        {
                cout << "How much (to the nearest dollar) did you"
                    << " spend on food during week " << week
                    << " ?:" << endl;

                cin >> dollars;

                total = total + dollars;
        }
        average = total / 4.0;

        cout << "Your weekly food bill over the chosen month is $"
            << average << endl << endl;
        do
        {
                cout << "Would you like to find the average for "
                    << "another month?";
```

continues

```
                cout << endl << "Enter Y or N" << endl;
                cin >> choice;
        } while(toupper(choice) != 'Y' && toupper(choice) != 'N');


    } while (toupper(choice) == 'Y');


    return 0;
}
```

This program prompts the user to input weekly food costs, to the nearest dollar (an integer) for a four-week period. The average weekly total for that month is output. Then the user is asked whether they want to repeat the calculation for a different month. The flow of this program is controlled by a do-while loop. The condition toupper(choice) == 'Y' allows the user to enter 'Y' or 'y' for yes. This makes the program more user friendly than if we just allowed 'Y'. Note the second do-while loop near the end of the program. This loop also utilizes toupper. Can you determine the purpose of this second loop? How would the execution of the program be affected if we removed this loop (but left in the lines between the curly brackets)?

String Constants

We have already talked about the character data type which includes letters, digits, and other special symbols such as $ and @. Often we need to put characters together to form strings. For example, the price "$1.99" and the phrase "one for the road!" are both strings of characters. The phrase contains blank space characters in addition to letters and an exclamation mark. In C++ a string is treated as a sequence of characters stored in consecutive memory locations. The end of the string in memory is marked by the null character \0. Do not confuse the null character with a sequence of two characters (i.e., \ and 0). The null character is actually an escape sequence. Its ASCII code is 0. For example, the phrase above is stored in computer memory as

o	n	e		f	o	r		t	h	e		r	o	a	d	!	\0

A **string constant** is a string enclosed in double quotation marks. For example,

"Learn C++"
"What time is it?"
"Code Word 7dF#c&Q"

are all string constants. When they are stored in the computer's memory, the null character is automatically appended. The string "Please enter a digit" is stored as

P	l	e	a	s	e		e	n	t	e	r		a		d	i	g	i	t	\0

When a string constant is used in C++, it is the memory address that is actually accessed. In the statement

```
cout << "Please enter a digit";
```

the memory address is passed to the cout object. cout then displays the consecutive characters until the null character is reached.

Storing Strings in Arrays

Often we need to access parts of a string rather than the whole string. For instance, we may want to alter characters in a string or even compare two strings. If this is the case, then a string constant is not what we need. Rather, a character array is the appropriate choice. When using character arrays, enough space to hold the null character must be allocated. For example:

```
char last[10];
```

This code defines a 10-element character array called `last`. However, this array can hold no more than 9 non-null characters since a space is reserved for the null character. Consider the following:

```
char last[10];
cout << "Please enter your last name using no more than 9 letters";
cin >> last;
```

If the user enters `Symon`, then the following will be the contents of the `last` array:

S	y	m	o	n	\0

Recall that the computer actually sees `last` as the beginning address of the array. There is a problem that can arise when using the `cin` object on a character array. `cin` does not "know" that `last` has only 10 elements. If the user enters `Newmanouskous` after the prompt, then `cin` will write past the end of the array. We can get around this problem by using the `getline` function. If we use

```
cin.getline(last,10)
```

then the computer knows that the maximum length of the string, including the null character, is 10. Consequently, `cin` will read until the user hits ENTER or until 9 characters have been read, whichever occurs first. Once the string is in the array, it can be processed character by character. In this next section we will see a program that uses `cin.getline()`.

Library Functions for Strings

The C++ library provides many functions for testing and manipulating strings. For example, to determine the length of a given string one can use the `strlen` function. The syntax is shown in the following code:

```
char line[40] = "A New Day";
int length;
length = strlen(line);
```

Here `strlen(line)` returns the length of the string including white spaces but not the null character at the end. So the value of `length` is 9. Note this is smaller than the size of the actual array holding the string.

To see why we even need a function such as `strlen`, consider the problem of reading in a string and then writing it backwards. If we only allowed strings of a fixed size, say length 29 for example, then the task would be easy. We simply read the string into an array of size 30 or more. Then write the 28th entry followed by the 27th entry and so on, until we reach the 0th entry. However, what if we wish to allow the user to input strings of different lengths? Now it is unclear where the end of the string is. Of course, we could search the array until we find

the null character and then figure out what position it is in. But this is precisely what the `strlen` function does for us. Sample Program 10.3 is a complete program that performs the desired task.

Sample Program 10.3:

```
#include <iostream>
#include <cstring>
using namespace std;

int main()
{
      char line[50];
      int length,count = 0;

      cout << "Enter a sentence of no more than 49 characters:\n";
      cin.getline(line,50);

      length = strlen(line);   // strlen returns the length of the
                               // string currently stored in line

      cout << "The sentence entered read backwards is:\n";

      for(count = length-1; count >= 0; count--)
      {
           cout << line[count];

      }

   cout << endl;
   return 0;
}
```

Sample Run 1:

Enter a sentence of no more than 49 characters:
luaP deiruB I
The sentence you entered printed backwards is:
I Buried Paul

Sample Run 2:

Enter a sentence of no more than 49 characters:
This sentence is too long to hold a mere 49 characters!
The sentence you entered printed backwards is:
arahc 94 erem a dloh ot gnol oot si ecnetnes sihT

Another useful function for strings is `strcat`, which concatenates two strings. `strcat(string1,string2)` attaches the contents of `string2` to the end of `string1`. The programmer must make sure that the array containing `string1` is large enough to hold the concatenation of the two strings plus the null character.

Consider the following code:

```
char string1[25] = "Total Eclipse ";    // note the space after the second
                                        // word - strcat does not insert a
                                        // space. The programmer must do this.
char string2[11] = "of the Sun";

cout << string1 << endl;
cout << string2 << endl;

strcat(string1,string2);

cout << string1 << endl;
```

These statements produce the following output:

```
Total Eclipse
of the Sun
Total Eclipse of the Sun
```

What would have happened if we had defined `string1` to be a character array
of size 20?

There are several other string functions such as `strcpy` (copies the second
string to the first string), `strcmp` (compares two strings to see if they are the
same or, if not, which string is alphabetically greater than the other), and `strstr`
(looks for the occurrence of a string inside of another string). Note that C-string
functions require the `cstring` header file. For more details on these string func-
tions and the others, see the text.

The `get` and `ignore` functions

There are several ways of inputting strings. We could use the standard >> extrac-
tion operator for a character array or string class object. However, we know that
using `cin` >> skips any leading whitespace (blanks, newlines). It will also stop
at the first trailing whitespace character. So, for example, the name "John Wayne"
cannot be read as a single string using `cin` >> because of a blank space between
the first and last names. We have already seen the `getline` function which does
allow blank spaces to be read and stored. In this section we will introduce the
`get` and `ignore` functions, which are also useful for string processing.

The `get` function reads in the next character in the input stream, including
whitespace. The syntax is

```
cin.get(ch);
```

Once this function call is made, the next character in the input stream is stored
in the variable `ch`. So if we want to input

```
$  X
```

we can use the following:

```
cin.get(firstChar);
cin.get(ch);
cin.get(secondChar);
```

where `firstChar`, `ch`, and `secondChar` are all character variables. Note that after the
second call to the `get` function, the blank character is stored in the variable `ch`.

The `get` function, like the `getline` function, can also be used to read strings. In this case we need two parameters:

```
cin.get(strName, numChar+1);
```

Here `strName` is a string variable and the integer expression `numChar+1` gives the number of characters that may be read into `strName`.

Both the `getline` and the `get` functions do not skip leading whitespace characters. The `get` statement above brings in the next input characters until it either has read `numChar+1` characters or it reaches the newline character `\n`. However, the newline character is not stored in `strName`. The null character is then appended to the end of the string. Since the newline character is not **consumed** (not read by the `get` function), it remains part of the input characters yet to be read.

Example:

```
char strName[21];
cin.get(strName,21);
```

Now suppose we input

John Wayne

Then "John Wayne" is stored in `strName`. Next input

My favorite westerns star John Wayne

In this case the string "My favorite westerns" is stored in `strName`.

We often work with records from a file that contain character data followed by numeric data. Look at the following data which has a name, hours worked, and pay rate for each record stored on a separate line.

Pay Roll Data

John Brown	7	12.50
Mary Lou Smith	12	15.70
Dominic DeFino	8	15.50

Since names often have imbedded blank spaces, we can use the `get` function to read them. We then use an integer variable to store the number of hours and a floating point variable to store the pay rate. At the end of each line is the '\n' character. Note that the end of line character is not consumed by reading the pay rate and, in fact, is the next character to be read when reading the second name from the file. This creates problems. Whenever we need to read through characters in the input stream without storing them, we can use the `ignore` function. This function has two arguments, the first is an integer expression and the second is a character expression. For example, the call

```
cin.ignore(80,'\n');
```

says to skip over the next 80 input characters but stop if a newline character is read. The newline character is consumed by the `ignore` function. This use of `ignore` is often employed to find the end of the current input line.

The following program will read the sample pay roll data from a file called payRoll.dat and show the result to the screen. Note that the input file must have names that are no longer than 15 characters and the first 15 positions of each line are reserved for the name. The numeric data must be after the 15th position in each line.

Sample Program 10.4:

```cpp
#include <fstream>
#include <iostream>
using namespace std;

const int MAXNAME = 15;

int main()
{
    ifstream inData;

    inData.open("payRoll.dat");
    char name[MAXNAME+1];
    int hoursWorked;
    float payRate;

    inData.get(name,MAXNAME+1);  // prime the read
    while (inData)
    {
      inData >> hoursWorked;
      inData >> payRate;

      cout << name << endl;
      cout << "Hours Worked " << hoursWorked << endl;
      cout << "Pay Rate " << payRate << " per hour"
           << endl << endl;

      inData.ignore(80,'\n');
        // This will ignore up to 80 characters but will
        // stop (ignoring) when it reads the \n which is
        // consumed.

      inData.get(name,MAXNAME+1);

    }

    return 0;
}
```

Summary of types of input for strings:

```
cin >> strName;          // skips leading whitespace. Stops at the first
                         // trailing whitespace (which is not consumed)


cin.get(strName, 21);    // does not skip leading whitespace
                         // stops when either 20 characters are read or
                         // '\n' is encountered (which is not consumed)


cin.ignore(200,'\n');    // ignores at most 200 characters but stops if
                         // newline (which is consumed) is encountered
```

Pointers and Strings

Pointers can be very useful for writing string processing functions. If one needs to process a certain string, the beginning address can be passed with a pointer variable. The length of the string does not even need to be known since the computer will start processing using the address and continue through the string until the null character is encountered.

Sample Program 10.5 below reads in a string of no more than 50 characters and then counts the number of letters, digits, and whitespace characters in the string. Notice the use of the pointer strPtr, which points to the string being processed. The three functions countLetters, countDigits, and countWhiteSpace all perform basically the same task—the while loop is executed until strPtr points to the null character marking the end of the string. In the countLetters function, characters are tested to see if they are letters. The if(isalpha(*strPtr)) statement determines if the character pointed at by strPtr is a letter. If so, then the counter occurs is incremented by one. After the character has been tested, strPtr is incremented by one to test the next character. The other two functions are analogous.

Sample Program 10.5:

```
#include <iostream>
#include <cctype>

using namespace std;

//function prototypes
int countLetters(char*);
int countDigits(char*);
int countWhiteSpace(char*);

int main()
{
    int numLetters, numDigits, numWhiteSpace;
    char inputString[51];


    cout <<"Enter a string of no more than 50 characters: "
         << endl << endl;
```

```
        cin.getline(inputString,51);

        numLetters = countLetters(inputString);
        numDigits = countDigits(inputString);
        numWhiteSpace = countWhiteSpace(inputString);

        cout << "The number of letters in the entered string is "
             << numLetters << endl;
        cout << "The number of digits in the entered string is "
             << numDigits << endl;
        cout << "The number of white spaces in the entered string is "
             << numWhiteSpace << endl;

        return 0;
}
//********************************************************************
//                    countLetters
//
// task:            This function counts the number of letters
//                  (both capital and lower case) in the string
// data in:         pointer that points to an array of characters
// data returned:   number of letters in the array of characters
//
//********************************************************************

int countLetters(char *strPtr)
{
    int occurs = 0;
    while(*strPtr != '\0')        // loop is executed as long as
                                  // the pointer strPtr does not point
                                  // to the null character which
                                  // marks the end of the string

    {
        if (isalpha(*strPtr))     // isalpha determines if
                                  // the character is a letter
            occurs++;
        strPtr++;
    }
    return occurs;
}
//********************************************************************
//                    countDigits
//
// task:            This function counts the number of digits
//                  in the string
// data in:         pointer that points to an array of characters
// data returned:   number of digits in the array of characters
//
//********************************************************************
```

continues

```
int countDigits(char *strPtr)
{
    int occurs = 0;
    while(*strPtr != '\0')
    {
        if (isdigit(*strPtr))   // isdigit determines if
                                // the character is a digit
            occurs++;
        strPtr++;
    }
    return occurs;
}


//********************************************************************
//                     countWhiteSpace
//
// task:              This function counts the number of whitespace
//                    characters in the string
// data in:           pointer that points to an array of characters
// data returned:     number of whitespaces in the array of
//                    characters
//
//********************************************************************

int countWhiteSpace(char *strPtr)   // this function counts the
                                    // number of whitespace characters.
                                    // These include, space, newline,
                                    // vertical tab, and tab
{
    int occurs = 0;
    while(*strPtr != '\0')
    {
        if (isspace(*strPtr))       // isspace determines if
                                    // the character is a
                                    // whitespace character
            occurs++;
        strPtr++;
    }
    return occurs;
}
```

PRE-LAB WRITING ASSIGNMENT

Fill-in-the-Blank Questions

1. The code cout << toupper('b'); causes a _____ to be displayed on the screen.

2. The data type returned by isalpha('g') is _____.

3. After the assignment statement `result = isdigit('$')`, `result` has the value _____.

4. The code `cout << tolower('#');` causes a _____ to be displayed on the screen.

5. The end of a string is marked in computer memory by the _____.

6. In `cin.getline(name,25)`, the 25 indicates that the user can input at most _____ characters into `name`.

7. Consider the following:

```
char message[35] = "Like tears in the rain";
int length;
length = strlen(message);
```

Then the value of `length` is _____.

8. Consider the code

```
char string1[30] = "In the Garden";
char string2[15] = "of Eden";
strcat(string1,string2);
cout << string1;
```

The output for this is _____.

9. The _____ header file must be included to access the `islower` and `isspace` character functions.

10. In C++, a string constant must be enclosed in _____ whereas a character constant must be enclosed in _____.

LESSON 10

LAB 10.1   Character Testing and String Validation

The American Equities investment company offers a wide range of investment opportunities ranging from mutual funds to bonds. Investors can check the value of their portfolio from the American Equities' web page. Information about personal portfolios is protected via encryption and can only be accessed using a password. The American Equities company requires that a password consist of 8 characters, 5 of which must be letters and the other 3 digits. The letters and digits can be arranged in any order. For example,

> **rt56AA7q**
> **123actyN**
> **1Lo0Dwa9**
> **myNUM741**

are all valid passwords. However, the following are all invalid:

> **the476NEw**   // It contains more than 8 characters (also more than 5
>   // letters)
> **be68moon**   // It contains less than 3 digits.
> **$retrn99**   // It contains only 2 digits and has an invalid character ('$')