

Folk Chapter 3

Secondary Storage and System
Software

Organization of Disks

- Disks *platters* rotate under a read/write head
- Disk partitions:
 - **Track** – *radial* partition
 - **Cylinder** – *radial* partition {a particular track from all platter surfaces}
 - **Sector** – *angular* partition
- **Read/Write Head**
 - At end of cantilever (actuator arm)
 - Moves radially

Organization of Disks

- To read a particular byte:
 - Appropriate surface, track, and sector identified by OS
 - Then *entire* sector read into buffer
 - Finally, desired byte located in buffer

Disk Capacities

- T – track, S – sector, b – byte, C – cylinder, D – drive
- $T = b/S * S/T = b/T$
- $C = b/T * T/C = b/C$
– $\{ = b/S * S/T * T/C \}$
- $D = b/C * C/D = b/D$
– $\{ = b/S * S/T * T/C * C/D \}$

Sector Organization of Tracks

- Logical vs. physical storage schema
 - Disk controller delay while processing sector can lead to not being ready for successive contiguous sector
 - Solved by interleaving
 - Obviated by sufficiently improved controller speeds in contemporary systems

Sector Organization of Tracks

- **Clusters**

- A fixed number of contiguous sectors
- Sector/Cluster correspondence maintained by the File Allocation Table (FAT)
- Sectors/cluster ratio adjustable by sysadmin

- **Extents**

- Clusters organized into a set of one or more extents.
- Each contiguous set of clusters for a file is an extent.
- Extents for a file are non-contiguous.
- Increasing the number of extents/file tends to increase the number of seeks/file

Sector Organization of Tracks

- **Fragmentation**

- Internal fragmentation w.r.t. sectors: wasted space in a sector when sector size is not an integral multiple of record size, and records not *spanning* two sectors.
- W.r.t clusters: wasted space in a cluster when the number of bytes/file isn't an integral multiple of the cluster size.

Block Organization of Tracks

- Disk blocks
 - Sizes can vary (some user control)
 - Not Unix system blocks
 - Alternative to sector organization
 - Obviates internal fragmentation problem
 - Blocking factor: records/block
 - Subblocks contain additional count and key info. regarding the block
 - Count: # of bytes in respective data block
 - Key: key for last record in data block – can allow more efficient searching by drive of track for block w. given key

Nondata Overhead

- Preformatting overhead
 - Sectors – at front of each sector: sector address, track address, condition (defective?); gaps and synchronization marks
 - Blocks – sub-block & inter-block gaps
 - Overhead can vary w. block sizes
 - More overhead w. blocks vs. sectors
 - Some overhead visible to programmer
 - Greater block sizes can lead to greater potential amount of internal track fragmentation

Costs of Disk Access

1. Seek Time (delay): time for move of r/w head to destination cylinder
 - Proportional to relative distance from starting cylinder to destination
 - Average seek distance: $1/3$ total # of tracks
2. Rotational Delay: time until destination sector/block is under r/w head
 - Can be all but eliminated for sequentially written files, both within a track and between tracks if the *beginning* of each successive track is staggered
3. Transfer Time (delay): time until all data (sectors/blocks) pass under r/w head

Issues Affecting Disk Performance

- Sequential access provides for much better r/w throughput than random access
- Block size affects performance
 - Larger blocks can more than linearly improve throughput, but at expense of fragmentation space. (table 3.2)
 - Dividing large blocks into smaller sub-blocks (e.g. 8 ½K blocks in one 4K block) maintains throughput w. less fragmentation by using sub-blocks for smaller files

Mitigating Disk as Bottleneck

- Disk performance lags well behind LAN performance. The following techniques help to mitigate the disk bottleneck:
 - **Multiprogramming/multiprocessing**: CPU attends to other programs/processes while waiting on disk I/O
 - **Disk *striping***: partition file onto several drives enabling *simultaneous* access, i.e. *parallelism*
 - RAID 0: Large blocks split into full tracks gather/scattered w. large disk controller cache (*buffer*)
 - **Disk *cache***: large block or RAM *mirroring pages* of data from a disk (*buffering*)

Tape

- Practical for sequential access only
- Used for archiving as *tertiary* storage
- Becoming increasingly obsolescent

CD-Rom

- Inexpensive, durable, high-capacity archival medium
- Write once – read many
- Very slow seeks
- Constant Linear Velocity (CLV) single track spiraling out from center
 - rotation speed proportional to radial placement of r/w head
 - Adds capacity at expense of seek time
 - Addressing by:
 - minutes (up to 70); seconds (60/min); sectors (75 2Kbytes/second)
 - Compare w. Constant Angular Velocity (CAV) of magnetic hard disks w. constant rotation speed

Storage Hierarchy

- Trade-off: capacity vs. access speed & throughput
 1. Primary
 1. Registers
 2. Level 1 cache
 3. level 2 [& 3] cache
 4. RAM
 2. Secondary
 1. Magnetic disks
 2. LAN
 3. Tertiary & Offline
 1. Removable media
 2. Broader networks (e.g. WWW)

Journey of an I/O Byte

- File manager layers of programs
 - Upper symbolic/logical file aspects
 - Opened? Type (e.g. binary)? Owner? Access permitted?
 - Lower physical layers
 - Info from FAT
- *System* I/O Buffer
 - Ensures that data organization in memory and disk respectively conform

Journey of an I/O Byte

- I/O Processor & Disk Controller
 - I/O *processor*: external processing device that *gather/scatters* byte groups to/from external devices offloading work from CPU
 - DMA (direct memory access): when the I/O processor can take data directly from RAM w.o. involving the CPU
 - Disk Controller: controls & monitors the disk.
 - Responds to queries & instructions from the I/O processor

Buffer Management

- Buffer Bottlenecks
 - Conflict between *input* and *output* function
 - Solved by separate input and output buffers
 - Defn.: **I/O bound** – CPU mostly idle *waiting* for I/O to be performed
 - Double buffering: alternating the roles of a pair of input and output buffers
 - Allows the OS to operate on one while the other's being loaded or emptied

Buffer Management

- Multiple buffering
 - Buffer pooling: buffer selected from pool of available upon demand
 - Replacement strategies
 - LRU (least recently used)
 - Best # of buffers system & problem dependent.
 - Copies between system and program buffers (*move mode*) can be eliminated if the system provides the program w. pointers to the system buffers (*locate mode*)
 - Scatter/Gather I/O: r/w with single instruction and multiple buffers – *scatter* input & *gather* output

Unix I/O

- Kernel: bottom layer of Unix OS. (Fig.3.23)
 - Views all I/O as byte sequences
 - No logical view of a file
 - Block (normal files), Character (term./printer), and Network (sockets) I/O each w. their own device & interface drivers
 - 4 tables:
 - File descriptor: points to entries in open file table
 - Open file table: file structure info. re. open files (ephemeral)
 - File allocation (*inode*) [persists as long as file exists]
 - Index nodes: hard link – file name to inode

File name linkage

- Hard link – file name in directory w. pointer to inode
 - Opening a file uses hard link to bring inode into memory & entry in open file table
 - Can be multiple hard links to inode, the number of which is maintained as an inode field – deleting a hard link decrements that count
- Soft (symbolic) link – an association between file names
 - Can leave *dangling* links post file deletion

Block I/O & Device Drivers

- Unix block: randomly addressable array of fixed blocks
- Device Driver: set of routines to perform I/O between a device and the an I/O buffer
 - Allows the kernel to view a device only abstractly

The Kernel and File Systems

- File systems reside on disk – components imported to memory by the kernel *as needed*
- File system and kernel are *separate* entities
 - File systems can be configured (*tuned*) to a specific device or usage pattern w.o. changing the kernel view of files
 - The kernel can operate with different & possibly multiple file systems